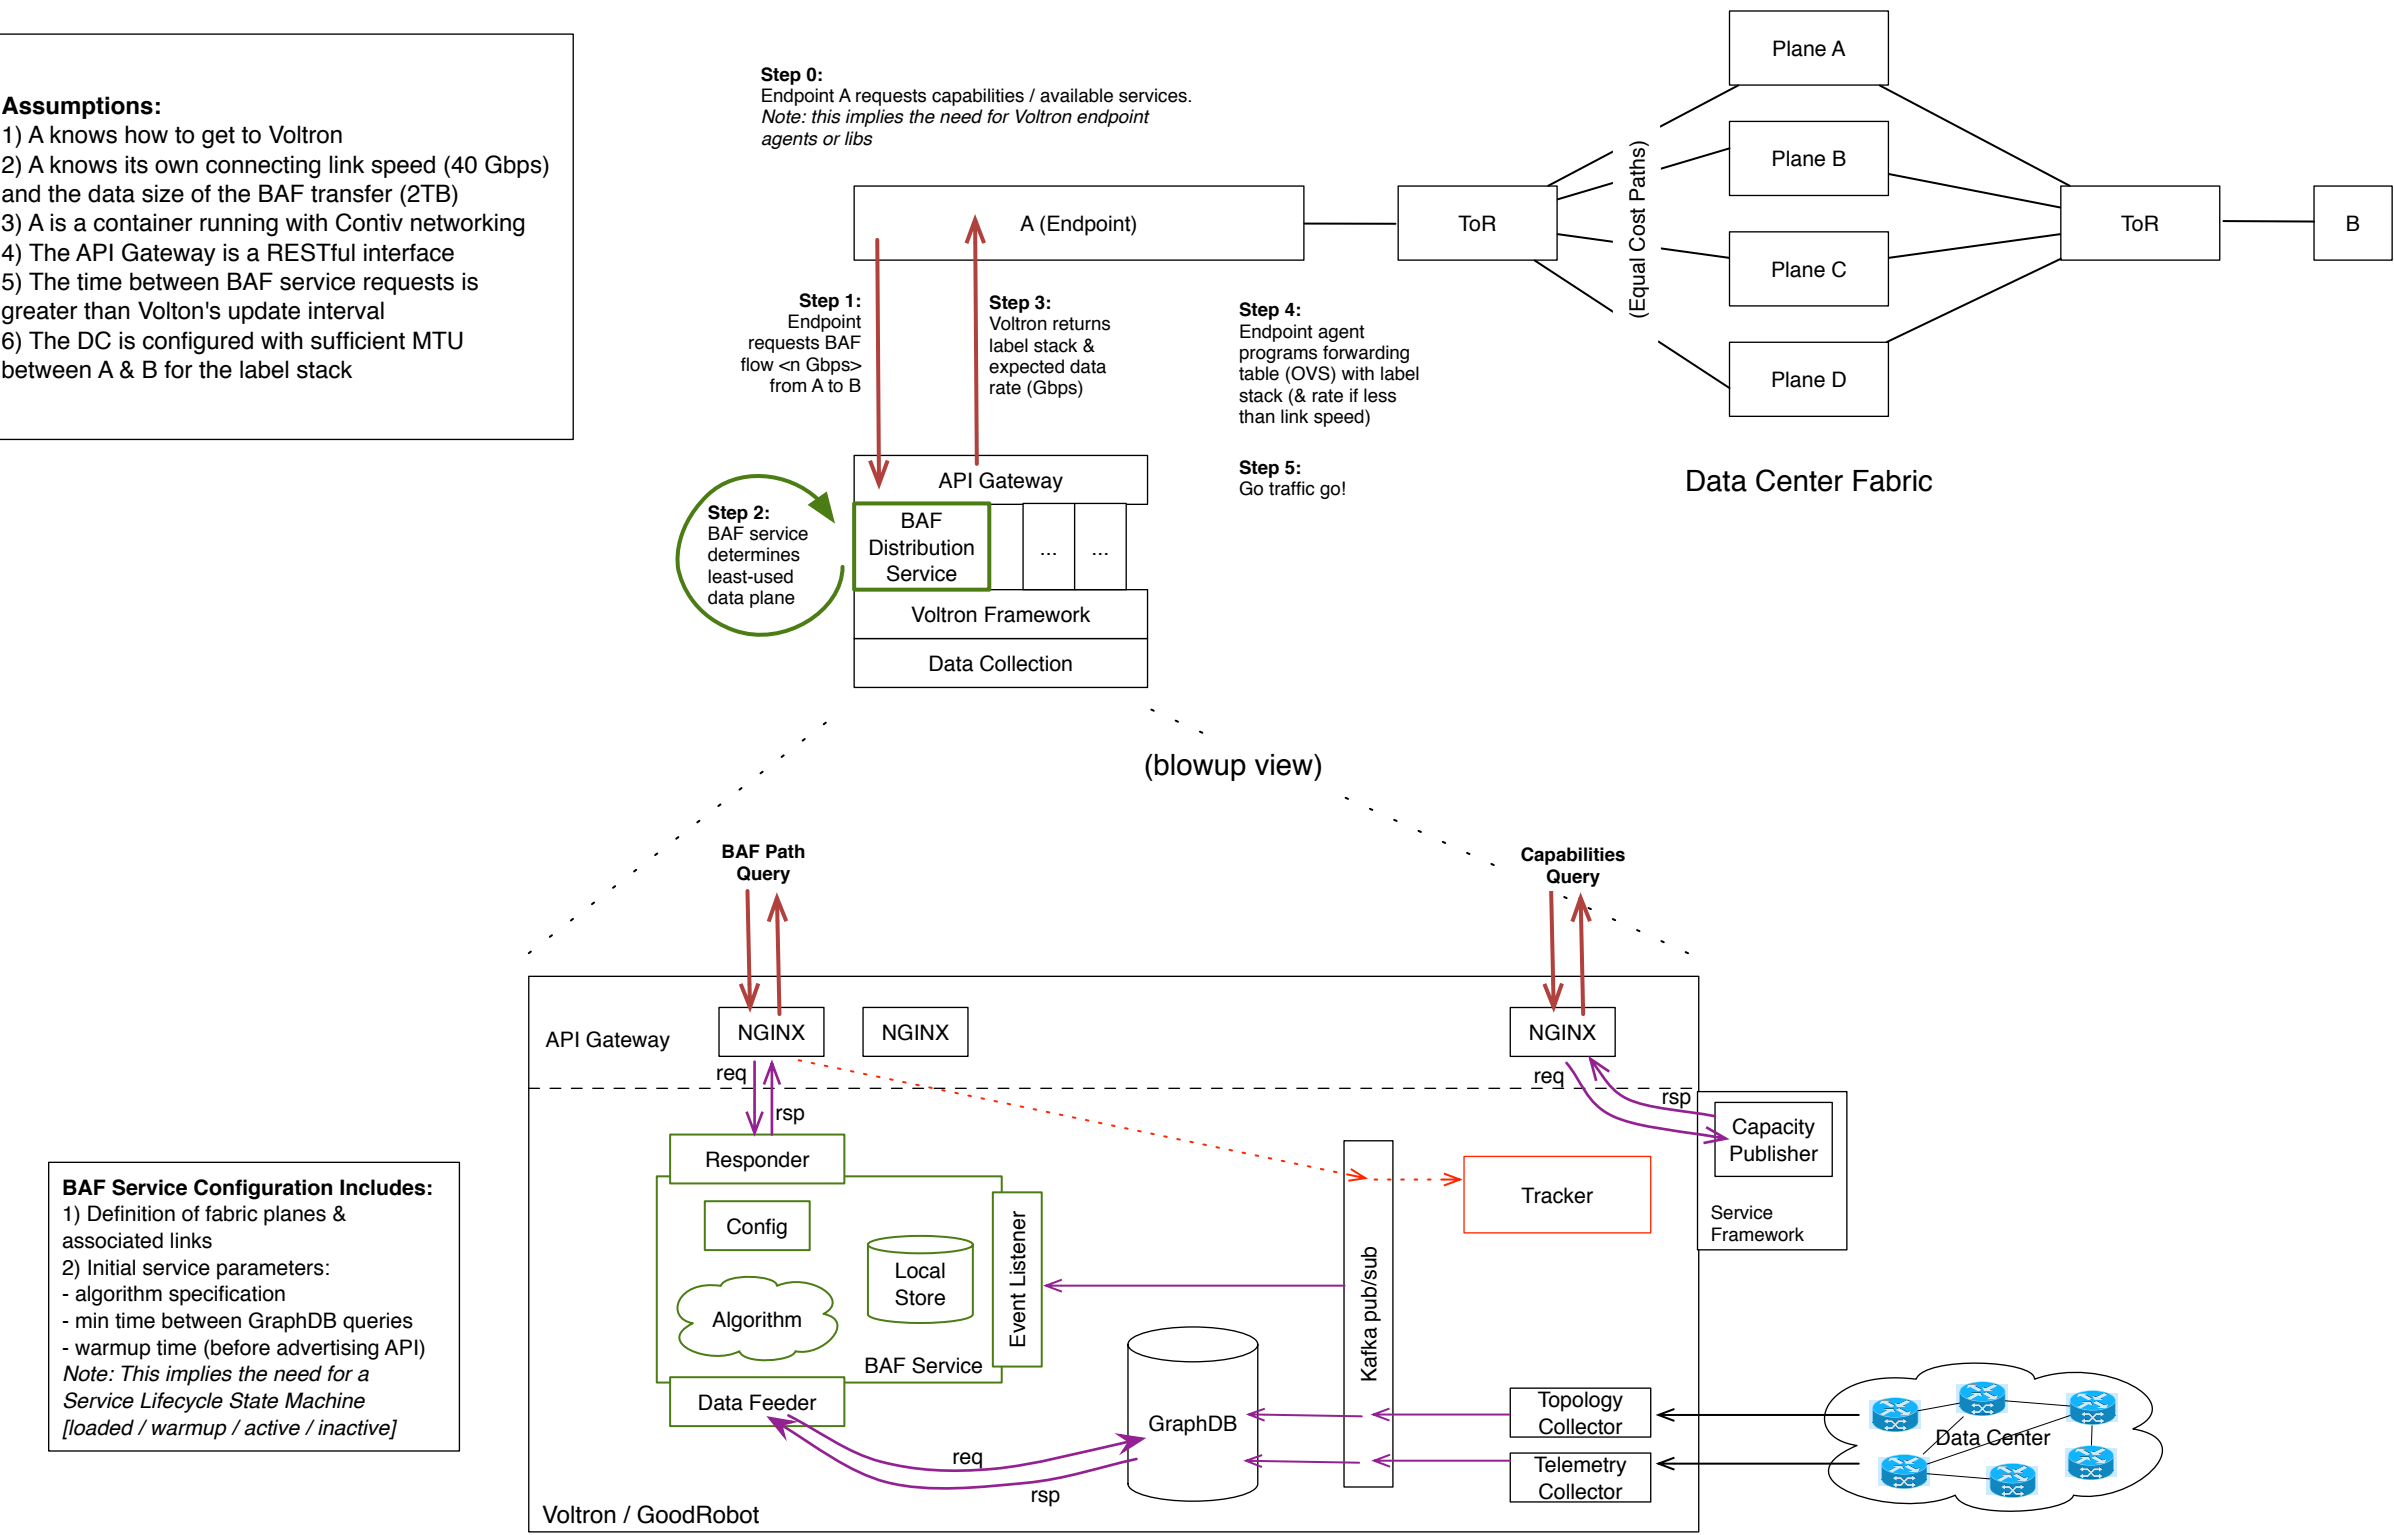


# Use Case: BAF, originating at A and going to B

Note: we are not intending to suggest paths for everybody. Just for special types of flows.

### Assumptions:

- 1) A knows how to get to Voltron
- 2) A knows its own connecting link speed (40 Gbps) and the data size of the BAF transfer (2TB)
- 3) A is a container running with Contiv networking
- 4) The API Gateway is a RESTful interface
- 5) The time between BAF service requests is greater than Volton's update interval
- 6) The DC is configured with sufficient MTU between A & B for the label stack



### BAF Service Configuration Includes:

- 1) Definition of fabric planes & associated links
- 2) Initial service parameters:
  - algorithm specification
  - min time between GraphDB queries
  - warmup time (before advertising API)*Note: This implies the need for a Service Lifecycle State Machine [loaded / warmup / active / inactive]*

### Situations to Consider

- \* How does BAF service deal with 100's to 1000's of endpoints? *one answer: have a set of best paths and round-robin through them*
- \* Due to telemetry data lag, need to be flexible of definition of "least used path" i.e. a set of low utilization paths
- \* Need to consider synchronization between "best path set" at time t, and new information received from telemetry
- \* Due to high bi-sectional bandwidth in the fabric, the most probable point of bottle necks is flow polarization to a single egress leaf node (note Max Flow Capacity of paths).
- \* Can we streamline service needs and data acquisition? Perhaps Kafka topics is a solution

\* Debug ability - what events & data do we need to log and/or track within Voltron?

## Use Case: Scaling the Fabric

Note: this use case describes a model wherein endpoint prefix scale in the containerized/virtualized data center exceeds forwarding table scale in the DC fabric itself. Initial case is single-tenant.

### Assumptions:

- 1) A knows how to get to Voltron
- 2) A is a container running with Contiv networking
- 3) A is a member of an application cluster and therefore cannot be NAT'd behind a host IP
- 4) A has an IPv4/v6 address that is not summarized at its local TOR
- 5) TOR prefix-SID is within the SRGB
- 6) Host prefix-SID is outside of the SRGB
- 7) The API Gateway is a RESTful interface
- 8) The DC is configured with sufficient MTU between A & B for the label stack

### Fabric Scale Service Configuration Includes:

- 1) Harvesting TOR prefix-SID and Host label data
    - Host label may simply be an EPE label for outgoing TOR interface
  - 2) Harvesting Host/container prefix data
  - 3) Initial service parameters:
    - creation of tuples from TOR-SID, Host-SID, prefix
    - min time between GraphDB queries
    - warmup time (before advertising API)
- Note: This implies the need for a Service Lifecycle State Machine [loaded / warmup / active / inactive]*

**Step 0:**  
Endpoint A requests capabilities / available services.  
*Note: this implies the need for Voltron endpoint agents or libs*

**Step 1:**  
Endpoint requests label stack to communicate with fellow app-cluster members

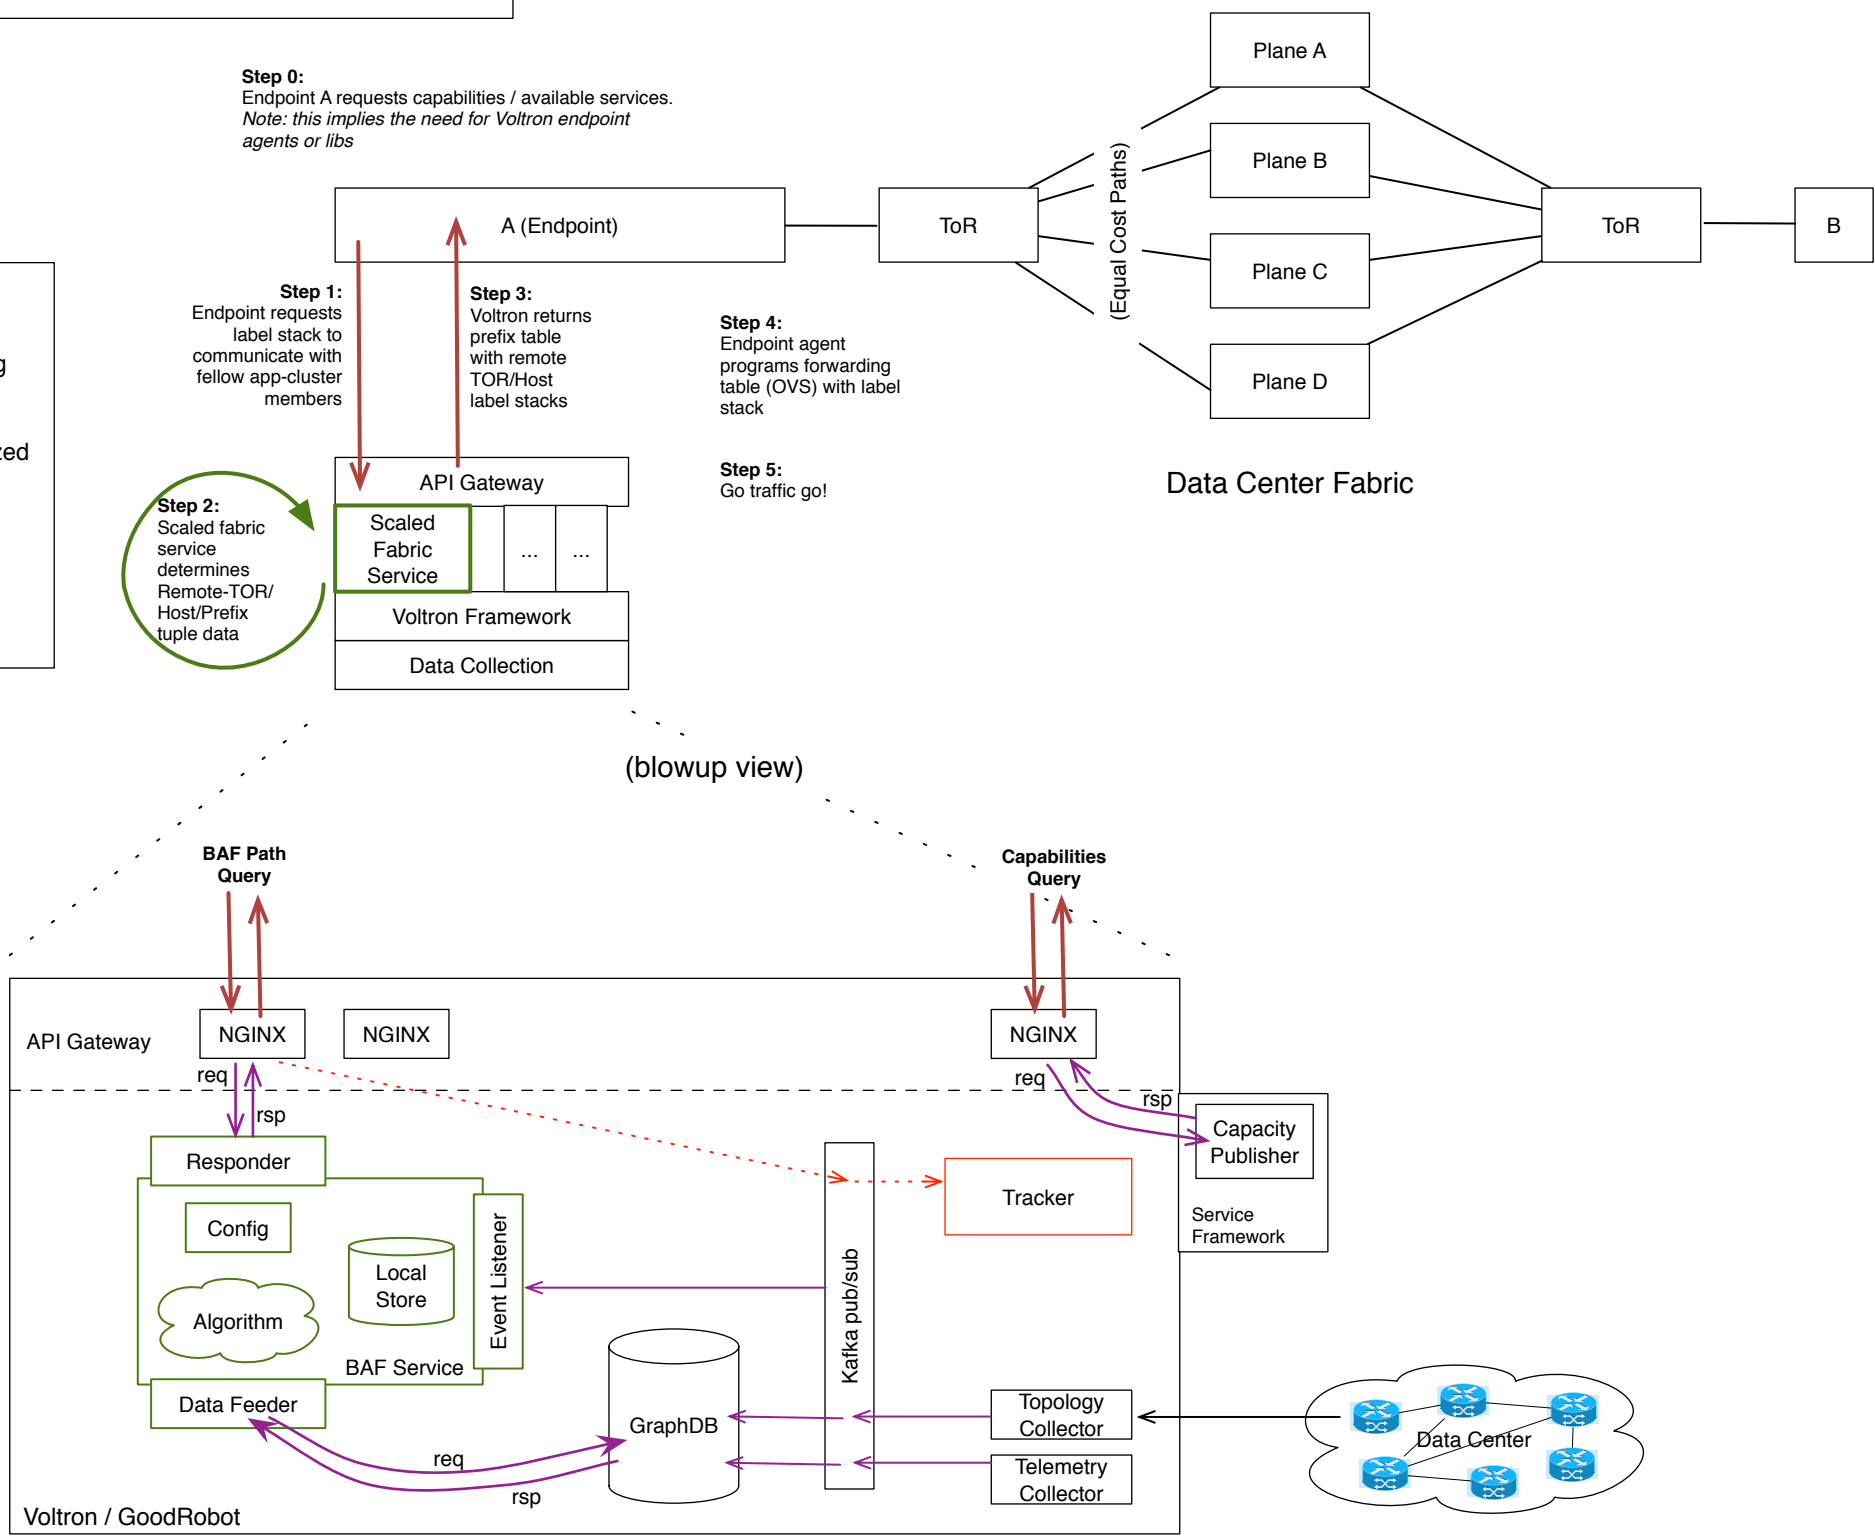
**Step 3:**  
Voltron returns prefix table with remote TOR/Host label stacks

**Step 4:**  
Endpoint agent programs forwarding table (OVS) with label stack

**Step 5:**  
Go traffic go!

**Step 2:**  
Scaled fabric service determines Remote-TOR/Host/Prefix tuple data

(blowup view)



### Situations to Consider

- \* How does BAF service deal with 100's to 1000's of endpoints? *one answer: have a set of best paths and round-robin through them*
- \* Due to telemetry data lag, need to be flexible of definition of "least used path" i.e. a set of low utilization paths
- \* Need to consider synchronization between "best path set" at time t, and new information received from telemetry
- \* Due to high bi-sectional bandwidth in the fabric, the most probable point of bottle necks is flow polarization to a single egress leaf node (note Max Flow Capacity of paths).
- \* Can we streamline service needs and data acquisition? Perhaps Kafka topics is a solution

\* Debug ability - what events & data do we need to log and/or track within Voltron?

## Use Case: Multi-Tenant Fabric

Note: this use case describes a multi-tenant containerized/virtualized data center wherein endpoints may be mobile within the DC and endpoints with different tenant IDs might have overlapping IPv4/v6 addresses

### Assumptions:

- 1) A knows how to get to Voltron
- 2) A is a container running with Contiv networking
- 3) Voltron is able to associate A's IPv4/v6 address with A's tenant\_ID
- 5) TOR prefix-SID is within the SRGB
- 6) Host prefix-SID is outside of the SRGB
- 7) The API Gateway is a RESTful interface
- 8) The DC is configured with sufficient MTU between A & B for the label stack

### Fabric Scale Service Configuration Includes:

- 1) Harvesting TOR prefix-SID and Host label data
    - Host label may simply be an EPE label for outgoing TOR interface
  - 2) Harvesting Container VPN/tenant, and prefix data
  - 3) Initial service parameters:
    - creation of tuples from TOR-SID, Host-SID, VPN-SID, prefix
    - min time between GraphDB queries
    - warmup time (before advertising API)
- Note: This implies the need for a Service Lifecycle State Machine  
[loaded / warmup / active / inactive]

**Step 0:**  
Endpoint A requests capabilities / available services.  
Note: this implies the need for Voltron endpoint agents or libs

**Step 1:**  
Endpoint requests label stack to communicate with fellow tenant members

**Step 2:**  
Multi-tenant fabric service determines Remote-TOR/Host/VPN/Prefix tuple data

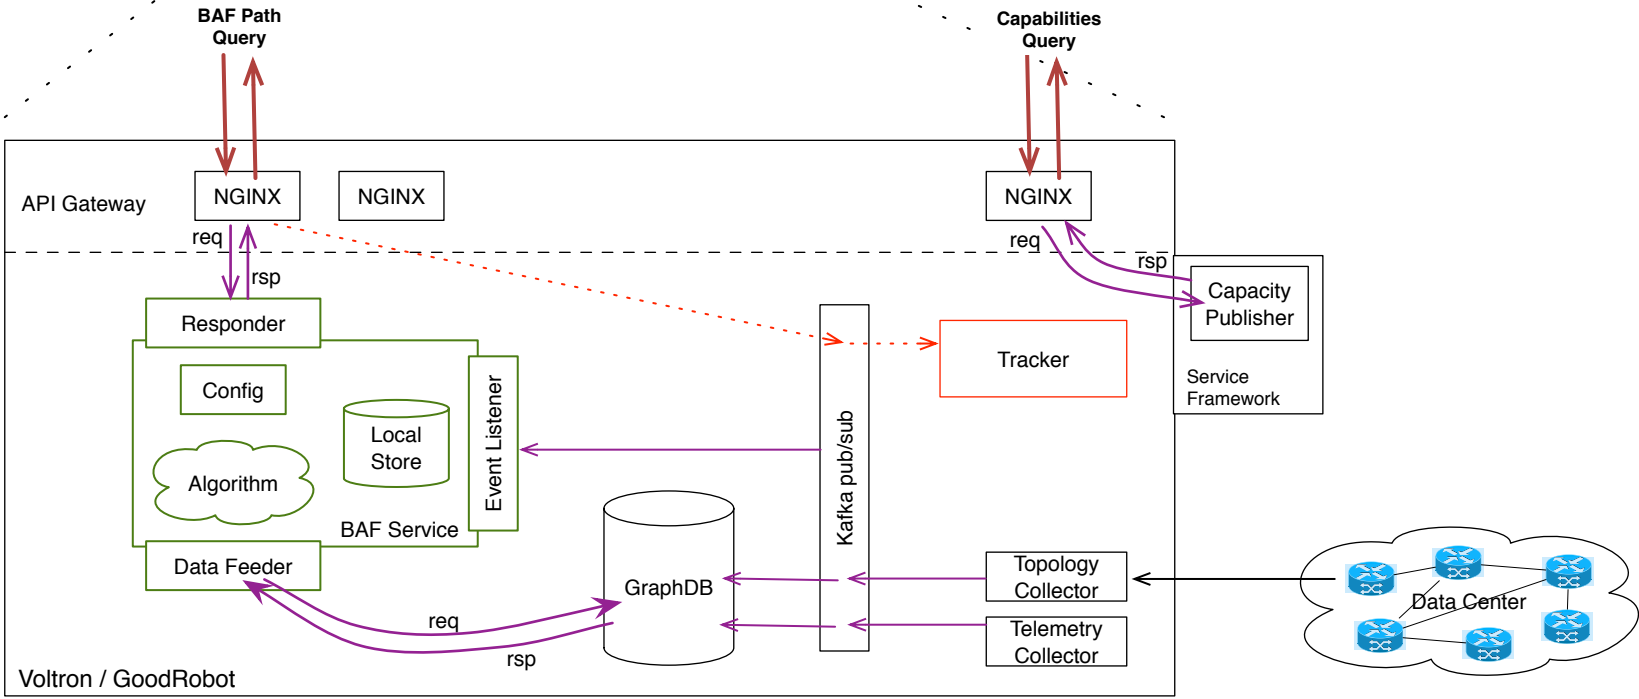
**Step 3:**  
Voltron returns prefix table with remote TOR/Host/VPN label stacks

**Step 4:**  
Endpoint agent programs forwarding table (OVS) with label stack

**Step 5:**  
Go traffic go!

Data Center Fabric

(blowup view)



### Situations to Consider

- \* How does BAF service deal with 100's to 1000's of endpoints? *one answer: have a set of best paths and round-robin through them*
- \* Due to telemetry data lag, need to be flexible of definition of "least used path" i.e. a set of low utilization paths
- \* Need to consider synchronization between "best path set" at time t, and new information received from telemetry
- \* Due to high bi-sectional bandwidth in the fabric, the most probable point of bottle necks is flow polarization to a single egress leaf node (note Max Flow Capacity of paths).
- \* Can we streamline service needs and data acquisition? Perhaps Kafka topics is a solution

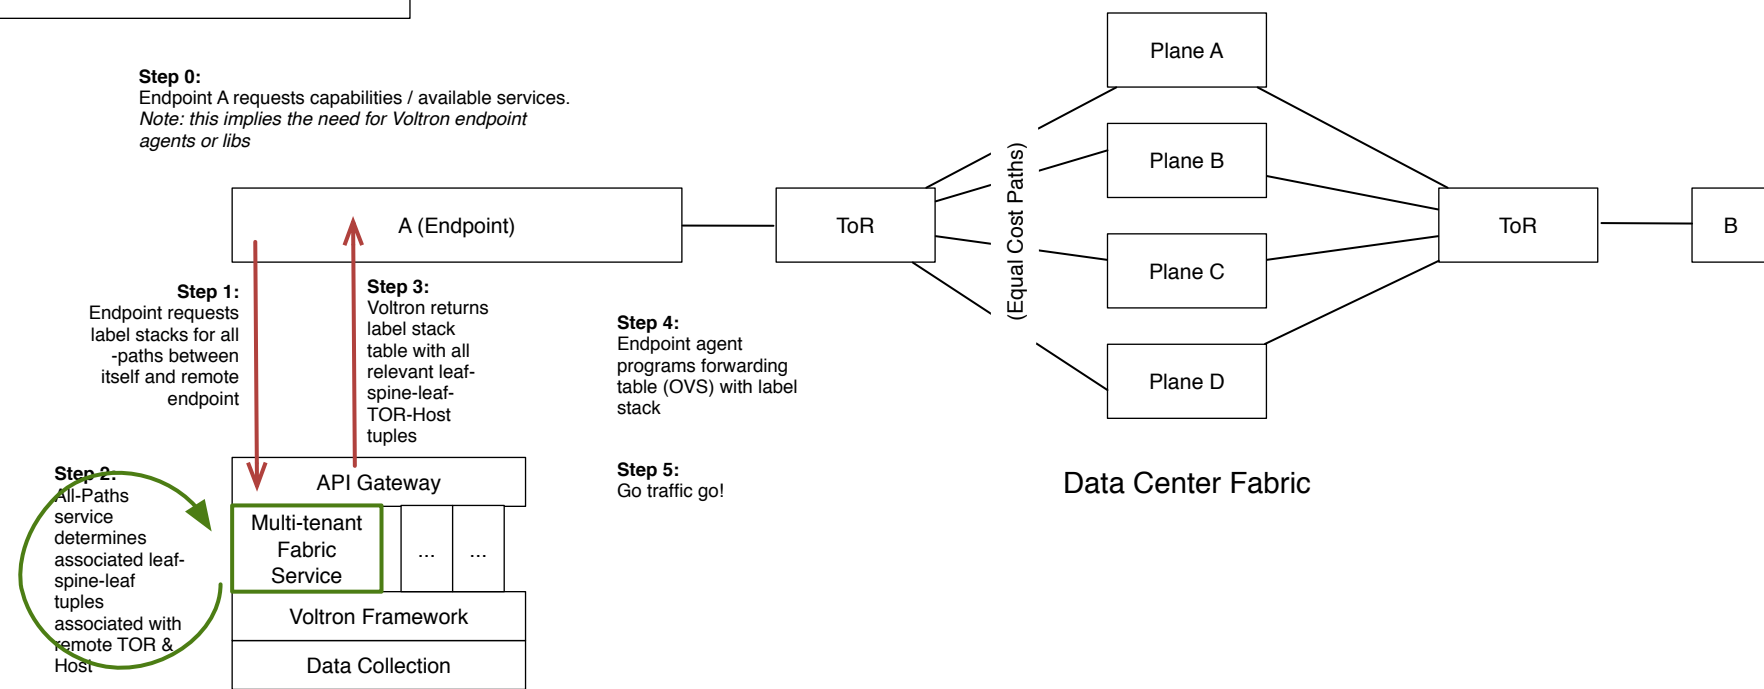
\* Debug ability - what events & data do we need to log and/or track within Voltron?

<h2>Use Case: All-Paths Service</h2>
<p>Note: this use case describes a scenario where an OAM tool needs to test the liveness (and perhaps latency) of all links in the fabric between itself and another host/endpoint. The micro service returns all-paths between the two endpoints in response to the request.</p>

<h2>Use Case: All-Paths Service</h2>
<p>Note: this use case describes a scenario where an OAM tool needs to test the liveness (and perhaps latency) of all links in the fabric between itself and another host/endpoint. The micro service returns all-paths between the two endpoints in response to the request.</p>

- 1) A knows how to get to Voltron
- 2) A is a container running with Contiv networking
- 3) All fabric nodes (Spine, leaf, TOR) have prefix-SIDs within the SRGB
- 4) Host prefix-SID is outside of the SRGB
- 7) The API Gateway is a RESTful interface
- 8) The DC is configured with sufficient MTU between A & B for the label stack

- 1) A knows how to get to Voltron
- 2) A is a container running with Contiv networking
- 3) All fabric nodes (Spine, leaf, TOR) have prefix-SIDs within the SRGB
- 4) Host prefix-SID is outside of the SRGB
- 7) The API Gateway is a RESTful interface
- 8) The DC is configured with sufficient MTU between A & B for the label stack



(blowup view)

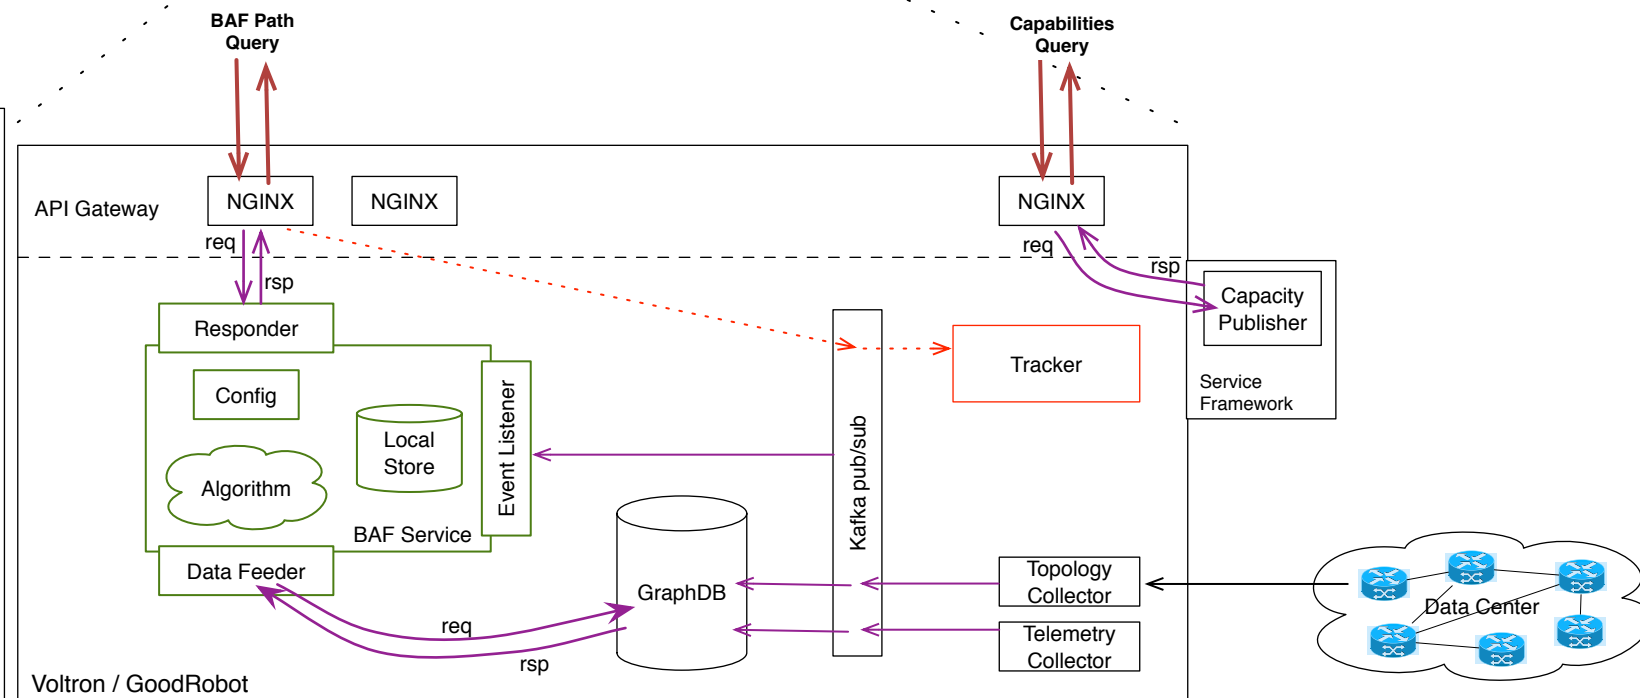
### All-Paths Service Configuration Includes:

- 1) Harvesting TOR prefix-SID and Host label data
  - Host label may simply be an EPE label for outgoing TOR interface
- 2) Harvesting of fabric node (spine and leaf) prefix-SID (and possibly adj-SID) data
- 3) Initial service parameters:
  - creation of tuples from leaf-spine-leaf to a given remote TOR/Host
  - example table size: assuming 3-stage CLOS built across 4-planes, TOR uplinks to 4 leaves (1 per plane), and leaf uplinks to 48 spines within its plane. Once traffic reaches a given leaf it has already entered a DC plane, therefore it has 48 paths to spine, then all spines have a single path to the appropriate egress leaf, so the table size is  $(4 \times 48 \times 1) = 192$  L-S-L tuples or paths.
  - warmup time (before advertising API)

*Note: This implies the need for a Service Lifecycle State Machine*

*[loaded / warmup / active / inactive]*

- ### All-Paths Service Configuration Includes:
- 1) Harvesting TOR prefix-SID and Host label data
    - Host label may simply be an EPE label for outgoing TOR interface
  - 2) Harvesting of fabric node (spine and leaf) prefix-SID (and possibly adj-SID) data
  - 3) Initial service parameters:
    - creation of tuples from leaf-spine-leaf to a given remote TOR/Host
    - example table size: assuming 3-stage CLOS built across 4-planes, TOR uplinks to 4 leaves (1 per plane), and leaf uplinks to 48 spines within its plane. Once traffic reaches a given leaf it has already entered a DC plane, therefore it has 48 paths to spine, then all spines have a single path to the appropriate egress leaf, so the table size is  $(4 \times 48 \times 1) = 192$  L-S-L tuples or paths.
    - warmup time (before advertising API)
- Note: This implies the need for a Service Lifecycle State Machine*
- [loaded / warmup / active / inactive]*



### Situations to Consider

- \* How does BAF service deal with 100's to 1000's of endpoints? *one answer: have a set of best paths and round-robin through them*
- \* Due to telemetry data lag, need to be flexible of definition of "least used path" i.e. a set of low utilization paths
- \* Need to consider synchronization between "best path set" at time t, and new information received from telemetry
- \* Due to high bi-sectional bandwidth in the fabric, the most probable point of bottle necks is flow polarization to a single egress leaf node (note Max Flow Capacity of paths).
- \* Can we streamline service needs and data acquisition? Perhaps Kafka topics is a solution

- \* Debug ability - what events & data do we need to log and/or track within Voltron?

- ### Situations to Consider
- \* How does BAF service deal with 100's to 1000's of endpoints? *one answer: have a set of best paths and round-robin through them*
  - \* Due to telemetry data lag, need to be flexible of definition of "least used path" i.e. a set of low utilization paths
  - \* Need to consider synchronization between "best path set" at time t, and new information received from telemetry
  - \* Due to high bi-sectional bandwidth in the fabric, the most probable point of bottle necks is flow polarization to a single egress leaf node (note Max Flow Capacity of paths).
  - \* Can we streamline service needs and data acquisition? Perhaps Kafka topics is a solution
- 
- \* Debug ability - what events & data do we need to log and/or track within Voltron?

### Situations to Consider

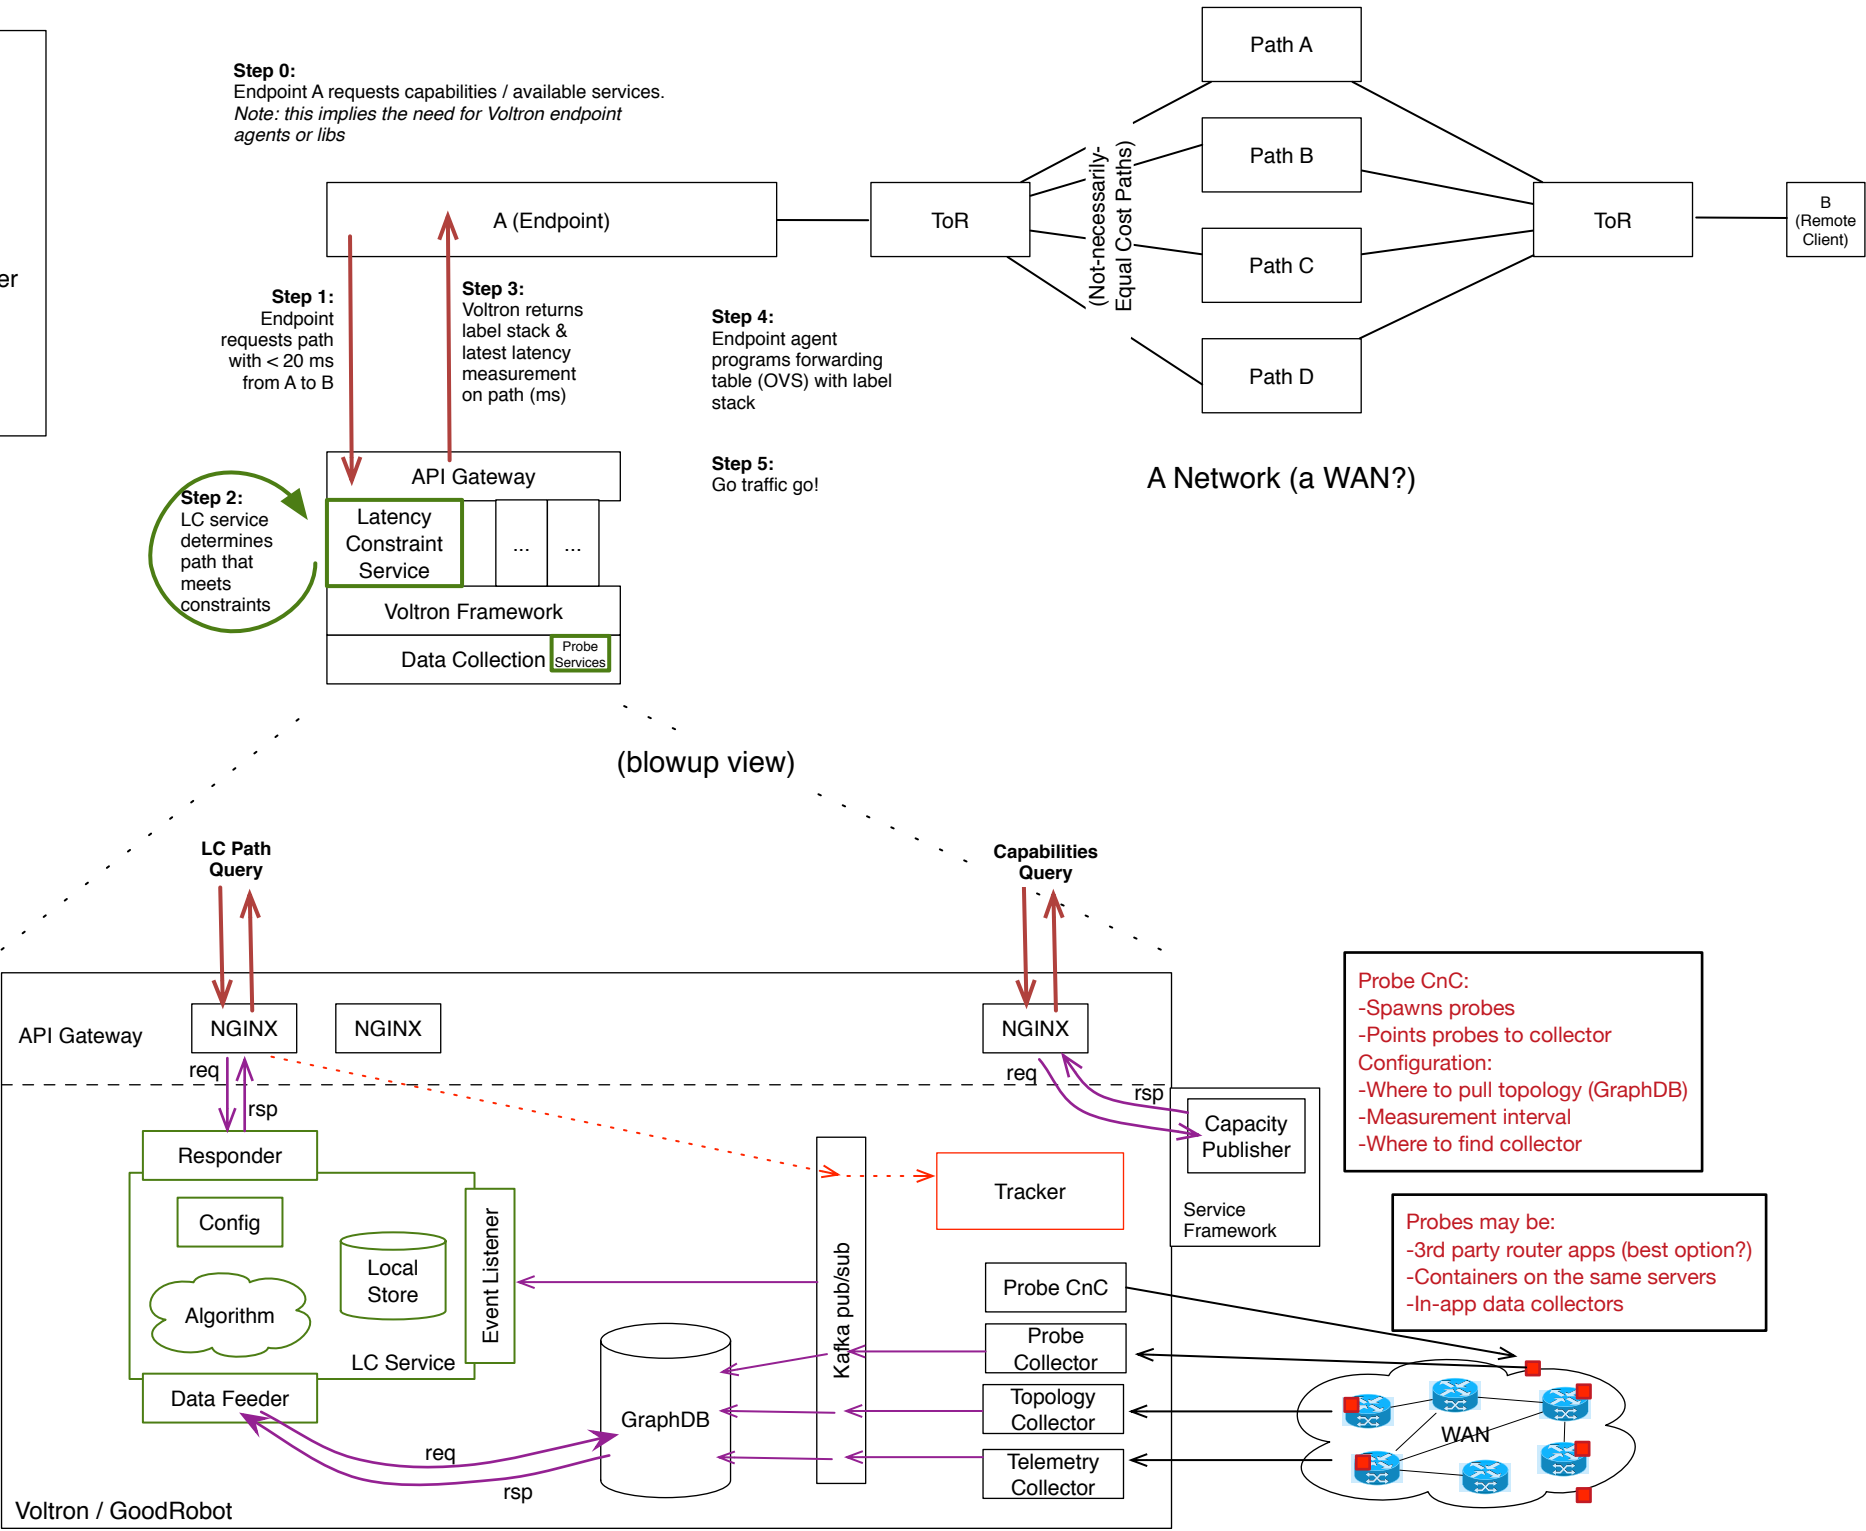
- \* How does BAF service deal with 100's to 1000's of endpoints? *one answer: have a set of best paths and round-robin through them*
- \* Due to telemetry data lag, need to be flexible of definition of "least used path" i.e. a set of low utilization paths
- \* Need to consider synchronization between "best path set" at time t, and new information received from telemetry
- \* Due to high bi-sectional bandwidth in the fabric, the most probable point of bottle necks is flow polarization to a single egress leaf node (note Max Flow Capacity of paths).
- \* Can we streamline service needs and data acquisition? Perhaps Kafka topics is a solution

- \* Debug ability - what events & data do we need to log and/or track within Voltron?

Use Case: Low Latency Flow (sub 20ms), originating at A and going to B

- Assumptions:**
- 1) A knows how to get to Voltron
  - 2) A knows its own connecting flow rate (1 Gbps) and its peak latency tolerance (20ms)
  - 3) A is a container running with Contiv networking
  - 4) The API Gateway is a RESTful interface
  - 5) The time between LC service requests is greater than Voltron's update interval (*required?*)
  - 6) The DC is configured with sufficient MTU between A & B for the label stack



**Situations to Consider**

- \* How does BAF service deal with 100's to 1000's of endpoints? *one answer: have a set of best paths and round-robin through them*
- \* Due to telemetry data lag, need to be flexible of definition of "least used path" i.e. a set of low utilization paths
- \* Need to consider synchronization between "best path set" at time t, and new information received from telemetry
- \* Due to high bi-sectional bandwidth in the fabric, the most probable point of bottle necks is flow polarization to a single egress leaf node (note Max Flow Capacity of paths).
- \* Can we streamline service needs and data acquisition? Perhaps Kafka topics is a solution

- \* Debug ability - what events & data do we need to log and/or track within Voltron?